


# Supervised Learning

- Would like to do **prediction**:  
estimate a function  $f(x)$  so that  $y = f(x)$
- Where  $y$  can be:
  - **Real number**: Regression
  - **Categorical**: Classification
  - **Complex object**:
    - Ranking of items, Parse tree, etc.
- **Data is labeled**:
  - Have many pairs  $\{(x, y)\}$ 
    - $x$  ... vector of real valued features
    - $y$  ... class ( $\{+1, -1\}$ , or a real number)



Training and  
test set

(1958)

F. Rosenblatt

**The perceptron: a probabilistic model  
for information storage and organization in the brain**  
*Psychological Review* 65:386-408

If we are eventually to understand the capability of higher organisms for perceptual recognition, generalization, recall, and thinking, we must first have answers to three fundamental questions:

1. How is information about the physical world sensed, or detected, by the biological system?
2. In what form is information stored, or remembered?
3. How does information contained in storage, or in memory, influence recognition and behavior?

The first of these questions is in the province of sensory physiology, and is

and the stored pattern. According to this hypothesis, if one understood the code or "wiring diagram" of the nervous system, one should, in principle, be able to discover exactly what an organism remembers by reconstructing the original sensory patterns from the "memory traces" which they have left, much as we might develop a photographic negative, or translate the pattern of electrical charges in the "memory" of a digital computer. This hypothesis is appealing in its simplicity and ready intelligibility, and a large family of theoretical brain models has been developed around the idea of a coded, representational memory (2, 3, 9, 14). The alternative

# Perceptron

# Linear models: Perceptron

## ■ Example: Spam filtering

	viagra	learning	the	dating	nigeria	<i>spam?</i>
$\vec{x}_1 = ($	1	0	1	0	0	$y_1 = 1$
$\vec{x}_2 = ($	0	1	1	0	0	$y_2 = -1$
$\vec{x}_3 = ($	0	0	0	0	1	$y_3 = 1$

- **Instance space  $x \in X$**  ( $|X| = n$  data points)
  - Binary feature vector  $x$  of word occurrences
  - $d$  features (words + other things,  $d \sim 100,000$ )
- **Class  $y \in Y$ :**
  - $y$ : Spam (+1), Ham (-1)


# Linear models for classification

- Binary classification:

$$f(x) = \begin{cases} 1 & \text{if } w_1 x_1 + w_2 x_2 + \dots + w_d x_d \geq \theta \\ 0 & \text{otherwise} \end{cases}$$

Decision  
boundary  
is linear

- **Input:** Vectors  $x_i$  and labels  $y_i$
- **Goal:** Find vector  $w = (w_1, w_2, \dots, w_d)$ 
  - Each  $w_i$  is a real number



Note:

$$\mathbf{x} \Leftrightarrow \langle \mathbf{x}, 1 \rangle \quad \forall \mathbf{x}$$

$$\mathbf{w} \Leftrightarrow \langle \mathbf{w}, -\theta \rangle$$

# Perceptron [Rosenblatt '57]

- **(very) Loose motivation: Neuron**

- Inputs are feature values

- Each feature has a weight  $w_i$




- Activation is the sum:

- $f(x) = \sum_i w_i x_i = w \cdot x$

- If the  $f(x)$  is:


  - Positive: predict +1

  - Negative: predict -1



# Perceptron: Estimating $w$

- **Perceptron:**  $y' = \text{sign}(w \cdot x)$
- **How to find parameters  $w$ ?**
  - Start with  $w_0 = 0$
  - Pick training examples  $x_t$  one by one (from disk)
  - Predict class of  $x$  using current weights
    - $y' = \text{sign}(w_t \cdot x_t)$
  - If  $y'$  is correct (i.e.,  $y_t = y'$ )
    - No change:  $w_{t+1} = w_t$
  - If  $y'$  is wrong: adjust  $w$ 
$$w_{t+1} = w_t + \eta \cdot y_t \cdot x_t$$
    - $\eta$  is the learning rate parameter
    - $x_t$  is the training example
    - $y_t$  is true class label ( $\{+1, -1\}$ )



# Perceptron Convergence

- **Perceptron Convergence Theorem:**
  - If there exist a set of weights that are consistent (i.e., the data is linearly separable) the perceptron learning algorithm will converge
- **How long would it take to converge?**
- **Perceptron Cycling Theorem:**
  - If the training data is not linearly separable the perceptron learning algorithm will eventually repeat the same set of weights and therefore enter an infinite loop
- **How to provide robustness, more expressivity?**

# Multiclass Perceptron


- **If more than 2 classes:**
  - Weight vector  $w_c$  for each class
    - Train one class vs. the rest
      - Example: 3-way classification  $y = \{A, B, C\}$
      - Train 3 classifiers:  $w_A$ : A vs. B,C;  $w_B$ : B vs. A,C;  $w_C$ : C vs. A,B

- **Calculate activation for each class**

$$f(x, c) = \sum_i w_{c,i} x_i = w_c \cdot x$$

- **Highest activation wins:**



$$c = \arg \max_c f(x, c)$$





# Issues with Perceptrons

- **Overfitting:**
- **Regularization:** if the data is not separable weights dance around
- **Mediocre generalization:**
  - Finds a “barely” separating solution



# Winnow Algorithm

## ■ Winnow algorithm

- Similar to perceptron, just different updates

**Initialize :**  $\theta = n$ ;  $w_i = 1$

**Prediction is 1 iff**  $w \bullet x \geq \theta$

**If no mistake : do nothing**

**If  $f(x) = 1$  but  $w \bullet x < \theta$  ,  $w_i \leftarrow 2w_i$  (if  $x_i = 1$ ) (promotion)**

**If  $f(x) = 0$  but  $w \bullet x \geq \theta$  ,  $w_i \leftarrow w_i/2$  (if  $x_i = 1$ ) (demotion)**

- Learns linear threshold functions

# Summary of Algorithms

Examples :  $\mathbf{x} \in \{0,1\}^n$ ;

Hypothesis :  $\mathbf{w} \in \mathbb{R}^n$

Prediction is 1 iff  $\mathbf{w} \bullet \mathbf{x} \geq \theta$

- **Additive** weight update algorithm  
[Perceptron, Rosenblatt, 1958]

$$\mathbf{w} \leftarrow \mathbf{w} + \eta_i y_j \mathbf{x}_j$$

If Class = 1 but  $\mathbf{w} \bullet \mathbf{x} \leq \theta$  ,  $w_i \leftarrow w_i + 1$  (if  $x_i = 1$ ) (promotion)

If Class = 0 but  $\mathbf{w} \bullet \mathbf{x} \geq \theta$  ,  $w_i \leftarrow w_i - 1$  (if  $x_i = 1$ ) (demotion)

- **Multiplicative** weight update algorithm  
[Winnow, Littlestone, 1988]

$$\mathbf{w} \leftarrow \mathbf{w} \eta_i \exp\{y_j \mathbf{x}_j\}$$

If Class = 1 but  $\mathbf{w} \bullet \mathbf{x} \leq \theta$  ,  $w_i \leftarrow 2w_i$  (if  $x_i = 1$ ) (promotion)

If Class = 0 but  $\mathbf{w} \bullet \mathbf{x} \geq \theta$  ,  $w_i \leftarrow w_i/2$  (if  $x_i = 1$ ) (demotion)

# Perceptron vs. Winnow

- **Perceptron**

- Online: can adjust to changing target, over time

- Advantages

- Simple
- Guaranteed to learn a linearly separable problem

- Limitations

- only linear separations
- only converges for linearly separable data
- not really “efficient with many features”

- **Winnow**

- Online: can adjust to changing target, over time

- Advantages

- Simple
- Guaranteed to learn a linearly separable problem
- **Suitable for problems with many irrelevant attributes**

- Limitations

- only linear separations
- only converges for linearly separable data
- not really “efficient with many features”